

MetaForge Control Console

Private AI Task Orchestration for R&D Teams

One-line pitch

MetaForge is a private AI control console that turns long-running engineering goals into auditable task flows with live status, evidence, and release gates.

The problem

R&D teams are adopting AI automation quickly, but they still struggle with:

- unclear task ownership and progress
- fragmented execution across chat, scripts, spreadsheets, and dashboards
- weak evidence trails for build, validation, and release decisions
- automation that works in demos but is too risky to operationalize
- failure recovery that depends on one person remembering the right sequence

The solution

MetaForge gives teams a single control plane for AI-assisted work:

- goal intake and task decomposition
- queueing and execution visibility
- evidence collection for build, smoke, and artifact checks
- release gating based on live facts, not stale state
- private deployment in a VM or internal environment

Who it is for

- Engineering teams using AI agents or automation in production workflows
- DevOps and platform teams that need traceable automation
- Research or prototyping groups that want a controlled path to operational use
- Security-conscious organizations that prefer private deployment

What it delivers

- less coordination overhead
- clearer execution status
- more trustworthy automation
- audit-ready evidence for delivery decisions
- fewer ad hoc scripts and fewer fragile handoffs

Product proof

- working browser demo: hosted URL TBD for external sharing
- screenshot proof: hosted image or public screen recording TBD
- internal demo artifact exists in the workspace and should not be linked in the external packet
- demo shows: goal intake -> queued tasks -> running tasks -> fresh evidence -> release gate

Traction snapshot

This is an honest pre-traction package. The current proof points are product and workflow evidence, not revenue claims:

- live control plane exists and runs privately
- evidence and release gates are implemented
- ToyOS regression proof and internal system validation exist
- design partner / pilot conversations are the next step

Team / why us

This section must be filled before formal investor outreach:

- founder bio: [insert]
- role split: [insert]
- prior shipped systems: [insert]
- relevant domain expertise: [insert]
- advisor / supporter quotes: [insert, if any]

Business model

- team subscription for private deployment
- enterprise subscription for internal rollout
- implementation and onboarding services for high-touch customers
- optional support / customization retainer

Fundraising ask

Draft structure for the external version. Replace the placeholders before sending:

- round: Seed
- instrument: [SAFE / priced equity]
- target raise: [amount TBD]
- valuation or cap: [TBD]
- use of funds:
 - product hardening and reliability
 - private deployment and enterprise features
 - design partner onboarding and pilots
 - go-to-market experiments and founder-led sales

Public demo materials

- public demo URL: [insert]
- public screenshot URL: [insert]
- short recording: [insert]

Why now

- teams want AI automation but need control
- private deployment and auditability are becoming table stakes
- simple “AI agent” demos do not survive real operational pressure

Current stage

- product narrative: ready
- working demo: ready
- deployment and integration packaging: in progress
- customer discovery: next step
- external fundraising package: draft pending founder inputs

Ask

Looking for pilot customers who want a controlled AI workflow layer for engineering operations, plus strategic partners who can help validate the private-deployment story. For investors, replace this line with the explicit funding ask, round, and use of funds above.